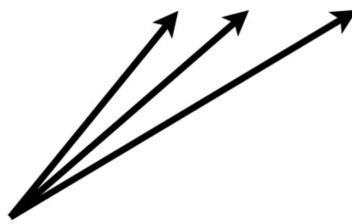
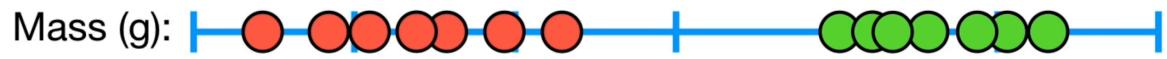
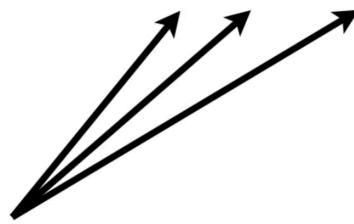
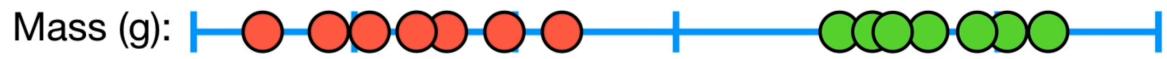


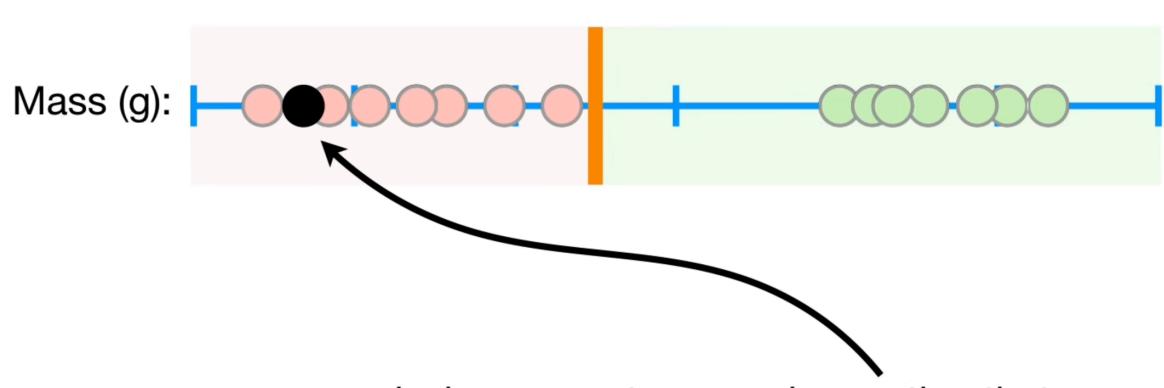
Let's start by imagining we measured
the mass of a bunch of mice...



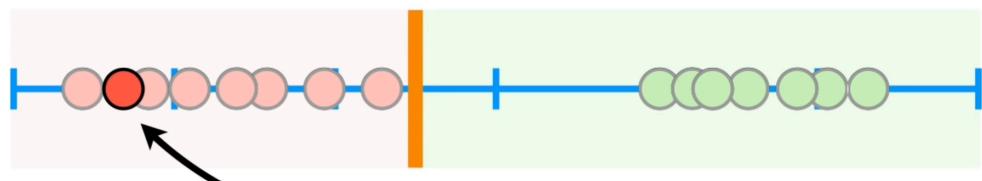
...and the **green dots** represent mice are **obese**.



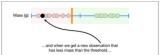
...and the **green dots** represent mice are **obese**.

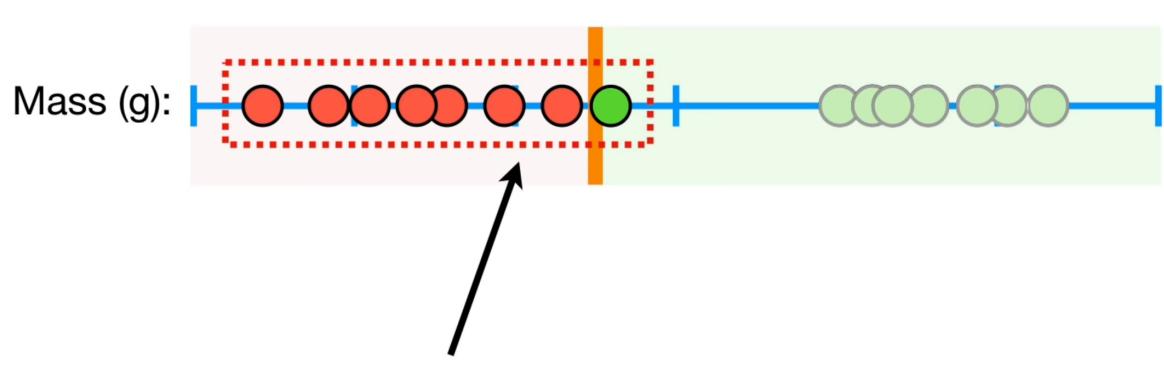


Mass (g):



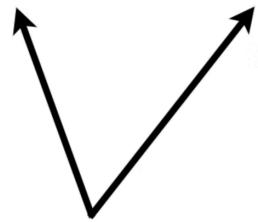
...we can classify it as ***not obese***.





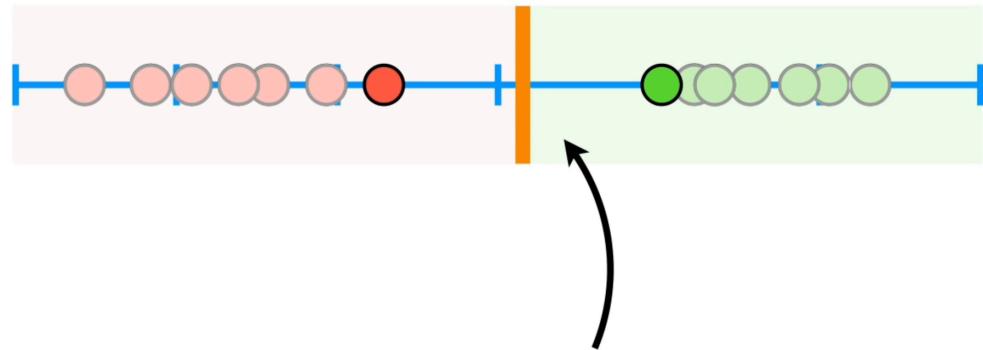
But that doesn't make sense, because it is much closer to the observations that are ***not obese***.

Mass (g):



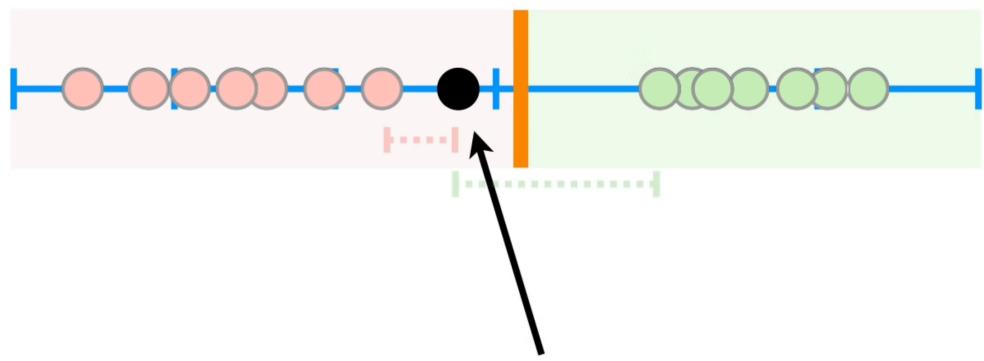
...we can focus on the observations on
the edges of each cluster...

Mass (g):

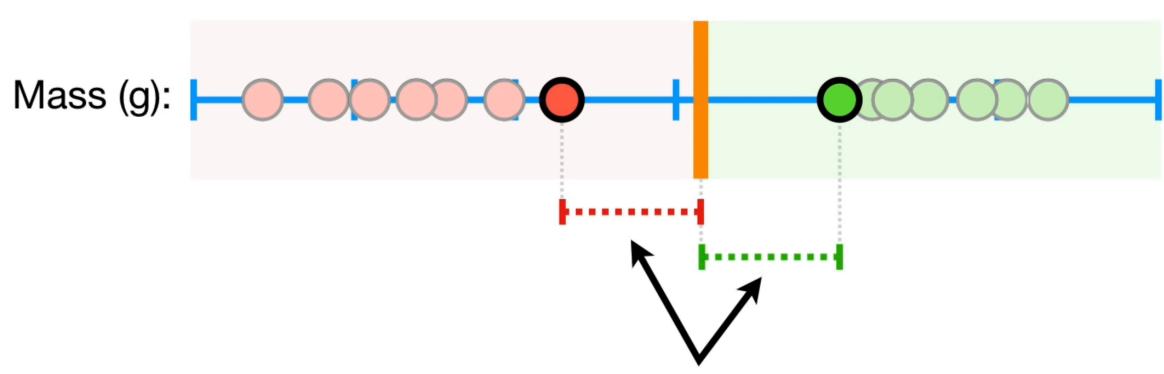


...and use the midpoint between
them as the threshold.

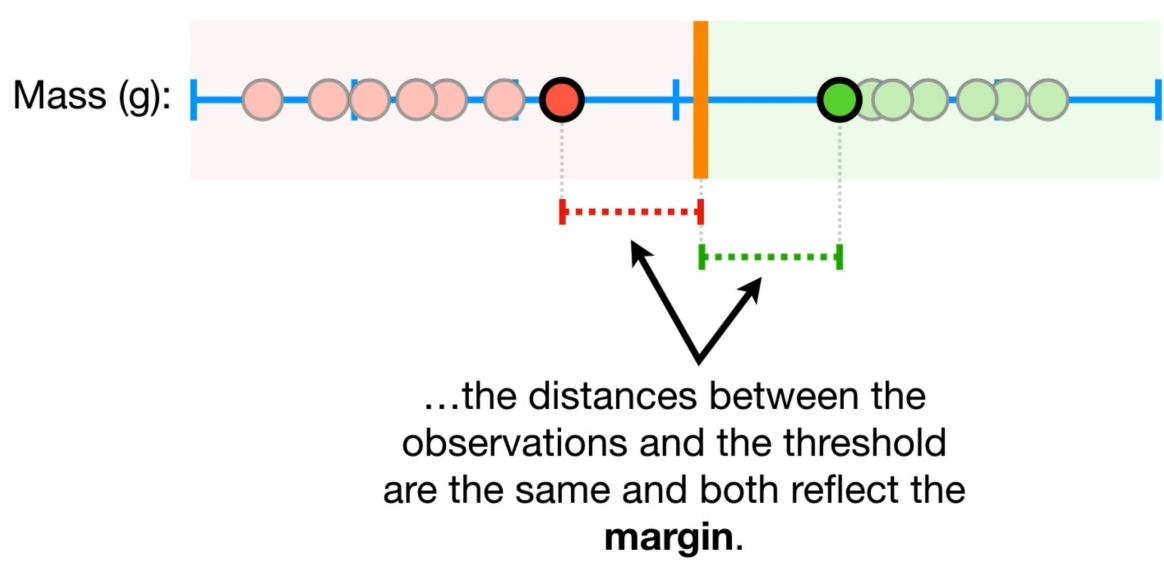
Mass (g):



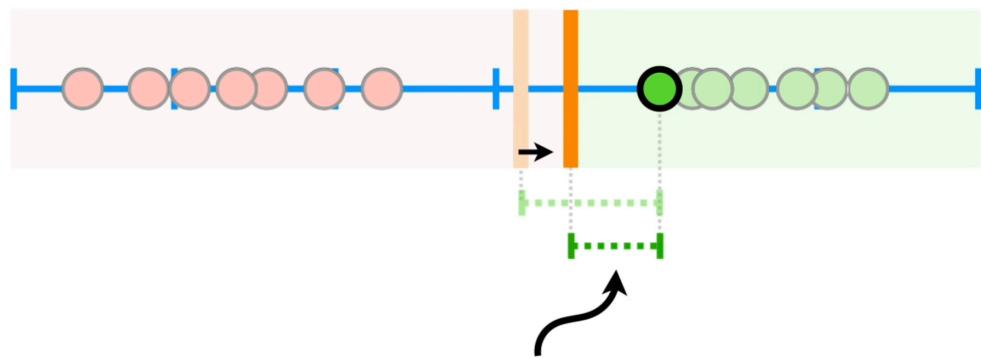
So it makes sense to classify this new observation as ***not obese***.



The shortest distance between
the observations and the
threshold is called the **margin**.

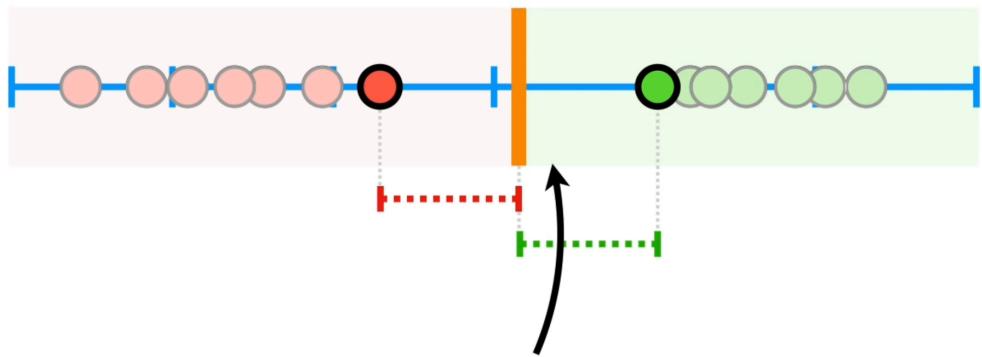


Mass (g):

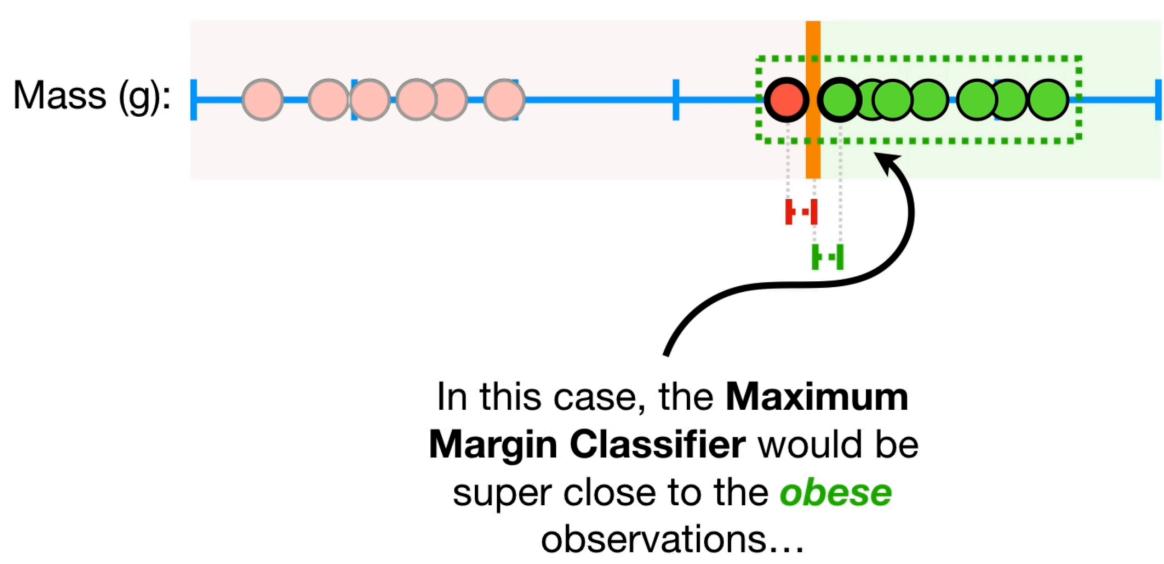


...and again, the **margin**
would be smaller.

Mass (g):

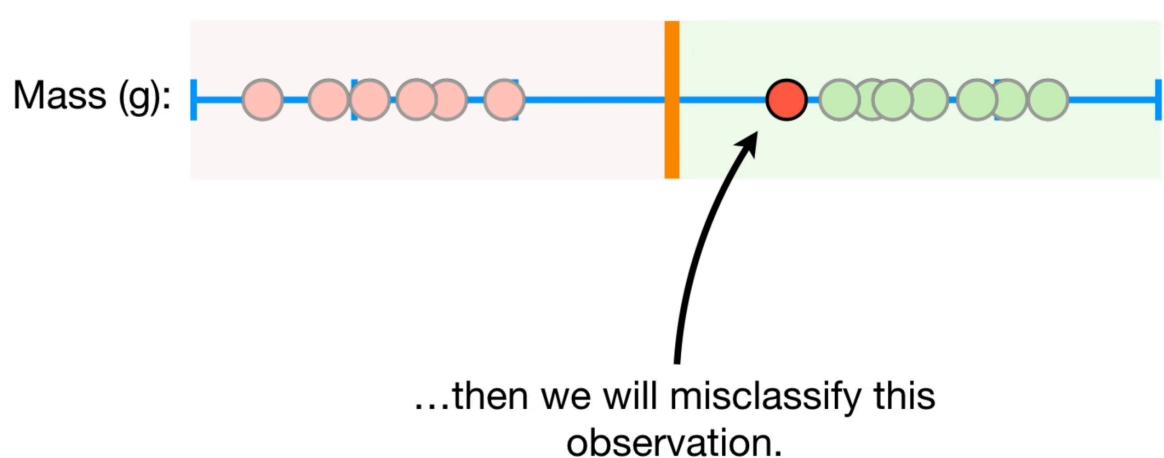


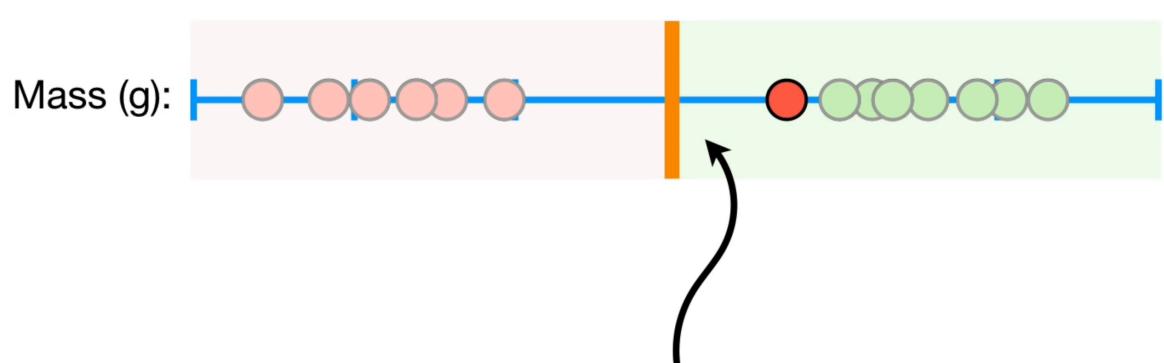
Maximal Margin Classifiers
seem pretty cool...



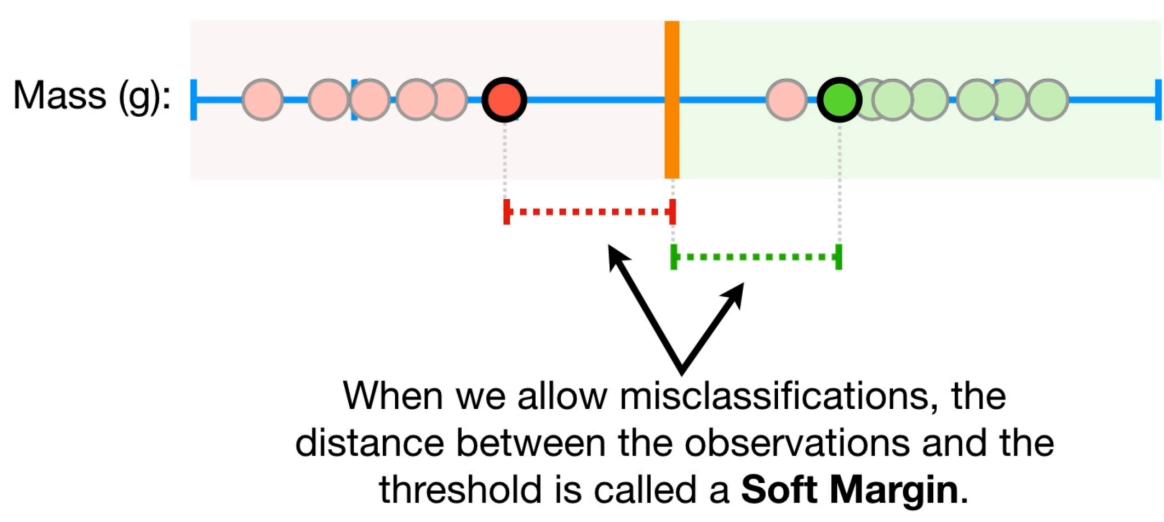


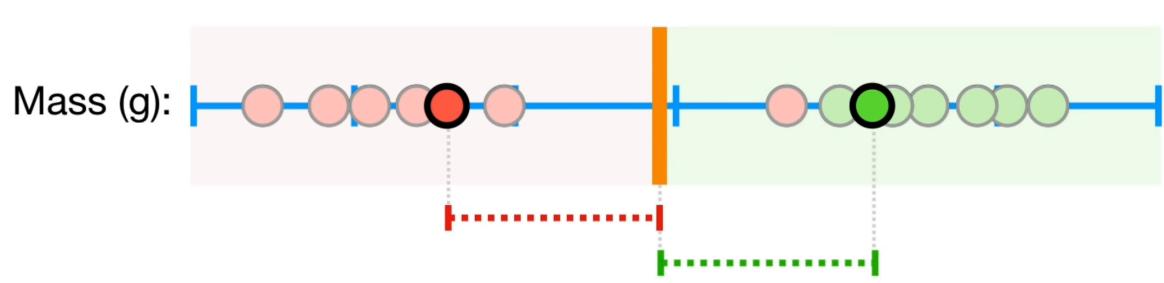
To make a threshold that is not so
sensitive to outliers we must **allow**
misclassifications.





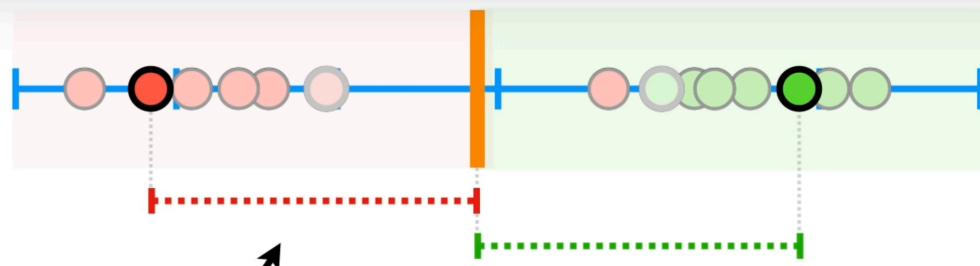
Choosing a threshold that allows misclassifications is an example of the **Bias/Variance Tradeoff** that plagues all of machine learning.





The answer is simple: We use **Cross Validation** to determine how many misclassifications and observations to allow inside of the **Soft Margin** to get the best classification.

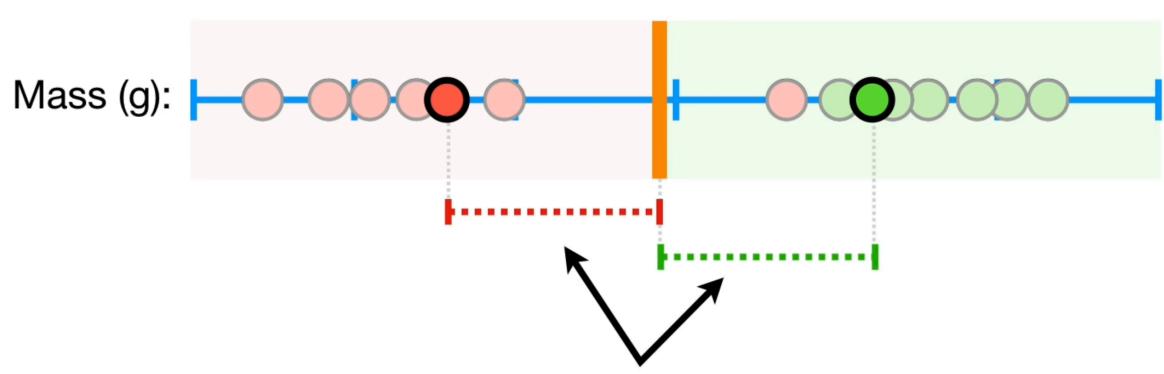
Mass (g):



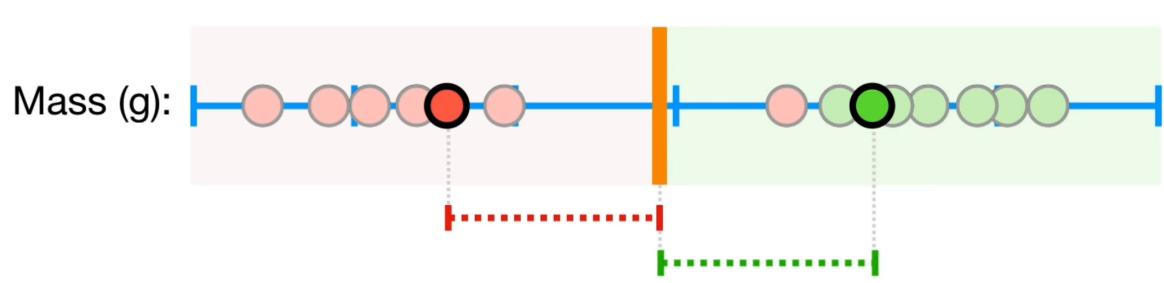
So the question is “How do we know
that this **soft margin**...

...is better than this **Soft Margin?**”

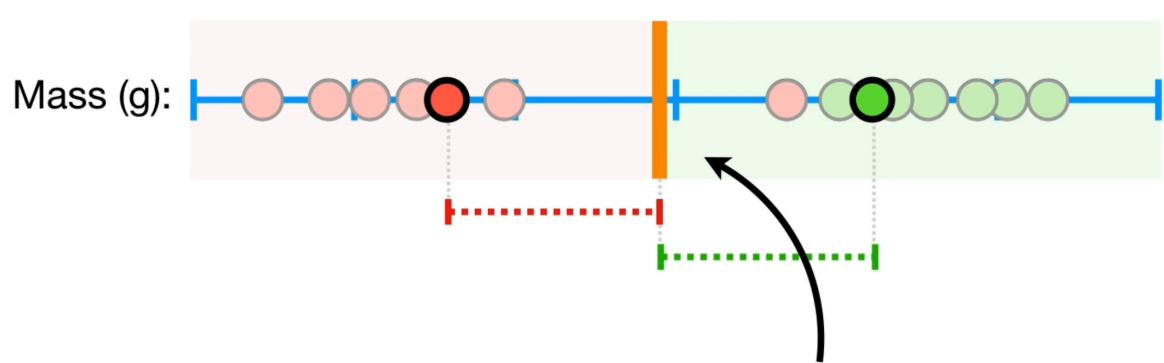




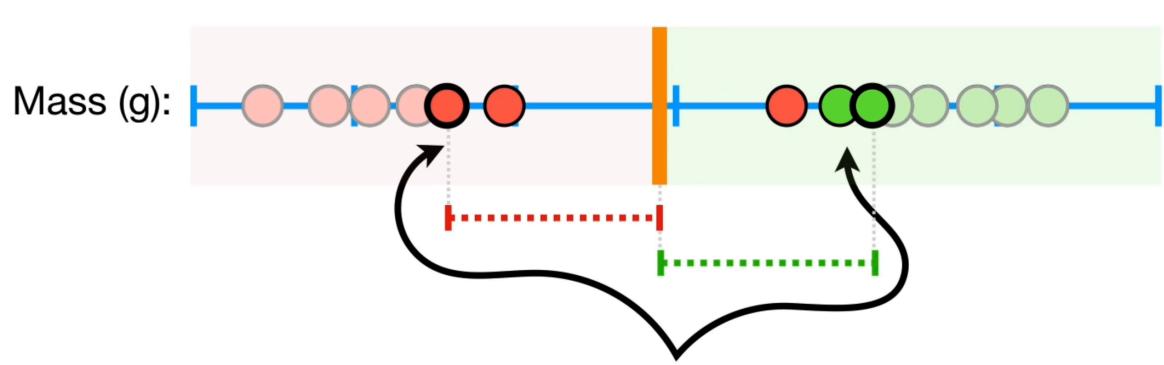
For example, if **Cross Validation** determined that this was the best **Soft Margin**...



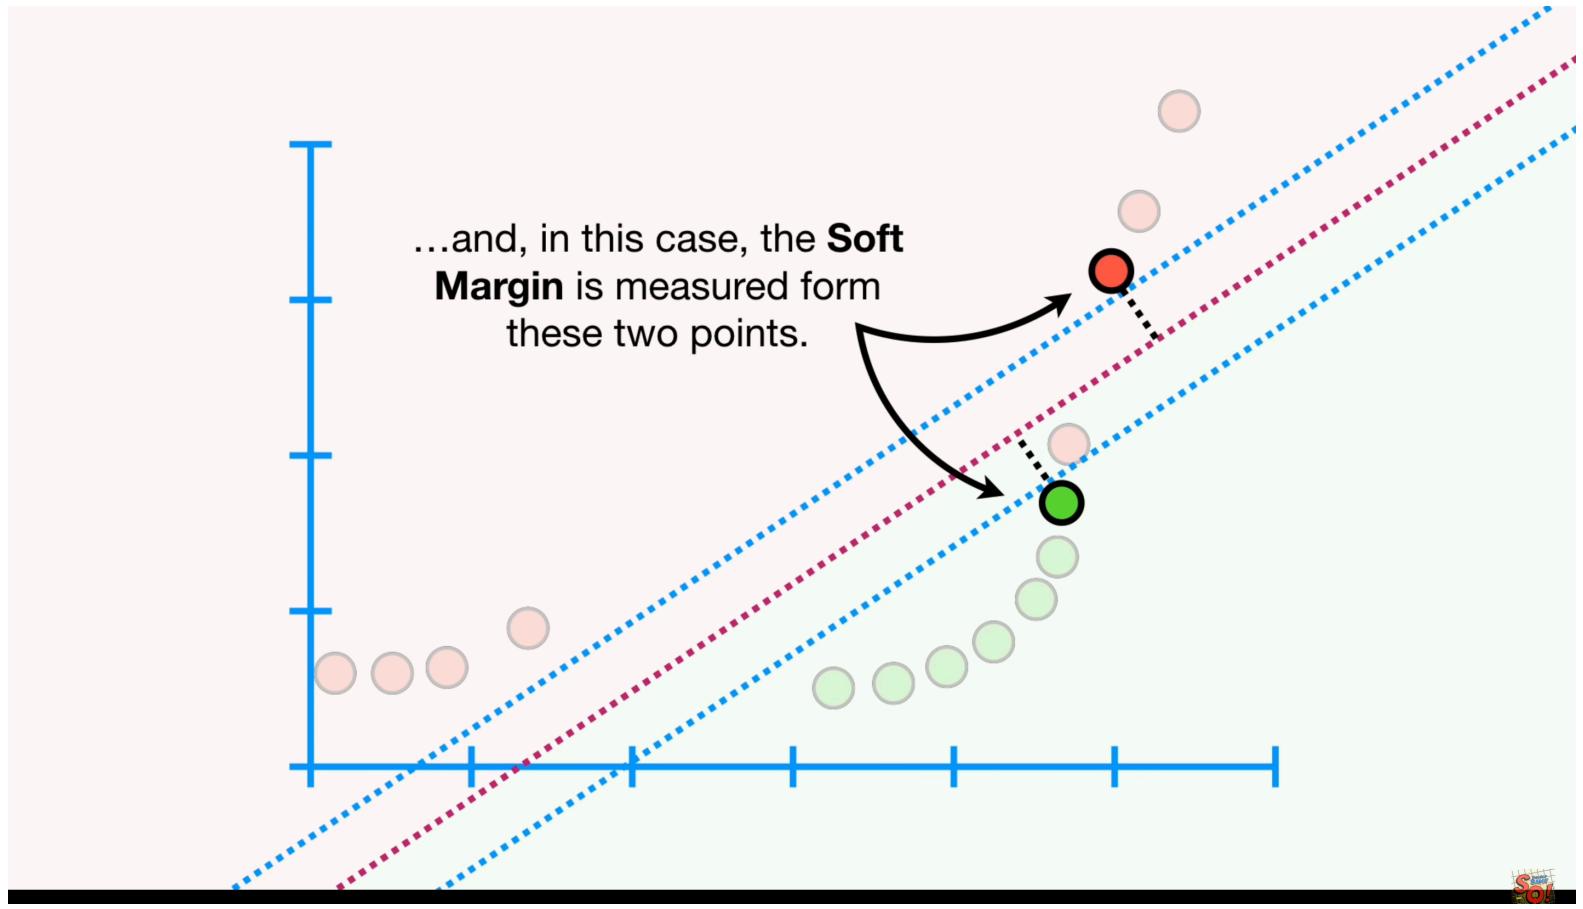
When we use a **Soft Margin** to determine the location of a threshold...



...then we are using a **Soft Margin Classifier** aka
a **Support Vector Classifier** to classify
observations.

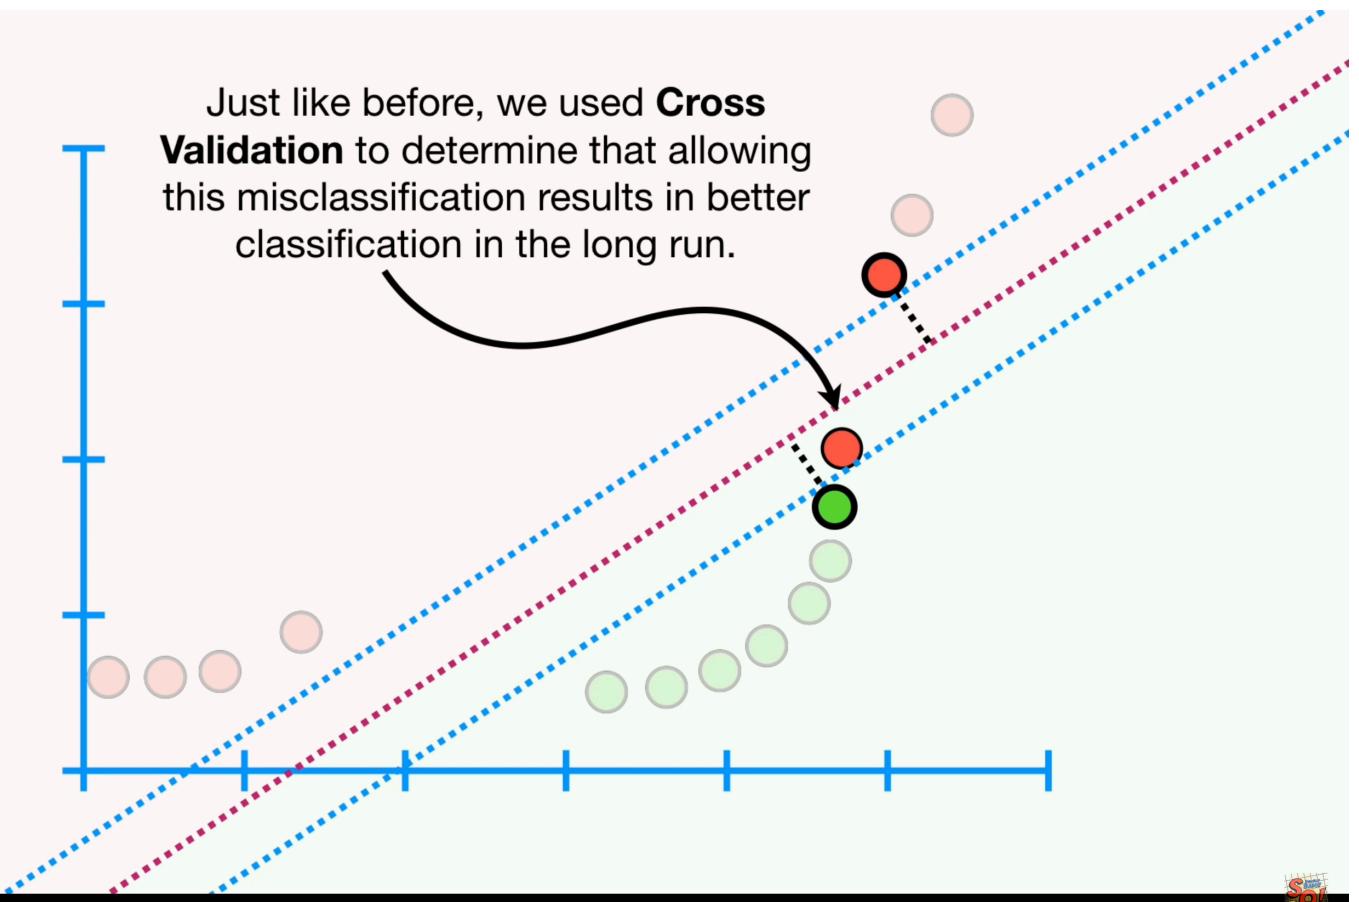


The name **Support Vector Classifier** comes from the fact that the observations on the edge *and within* the **Soft Margin** are called **Support Vectors**.

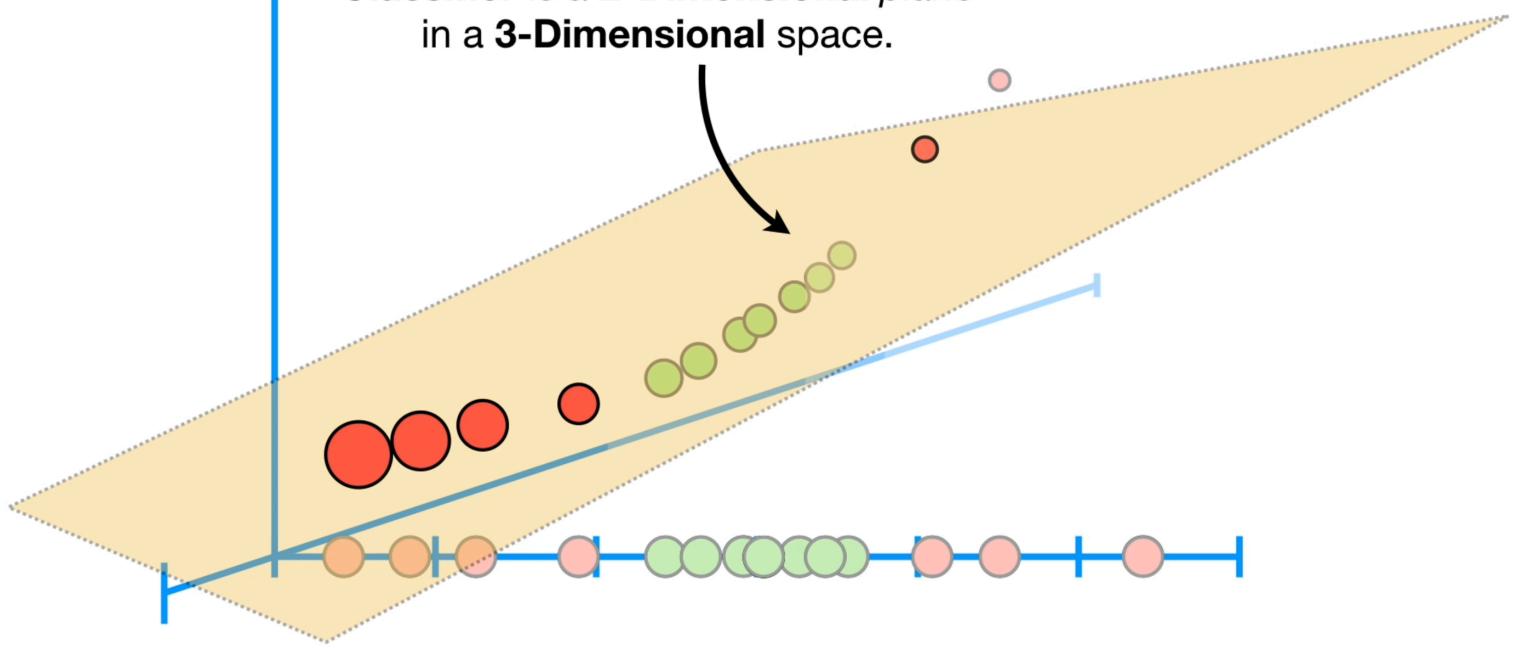


...and, in this case, the **Soft Margin** is measured from these two points.

Just like before, we used **Cross Validation** to determine that allowing this misclassification results in better classification in the long run.



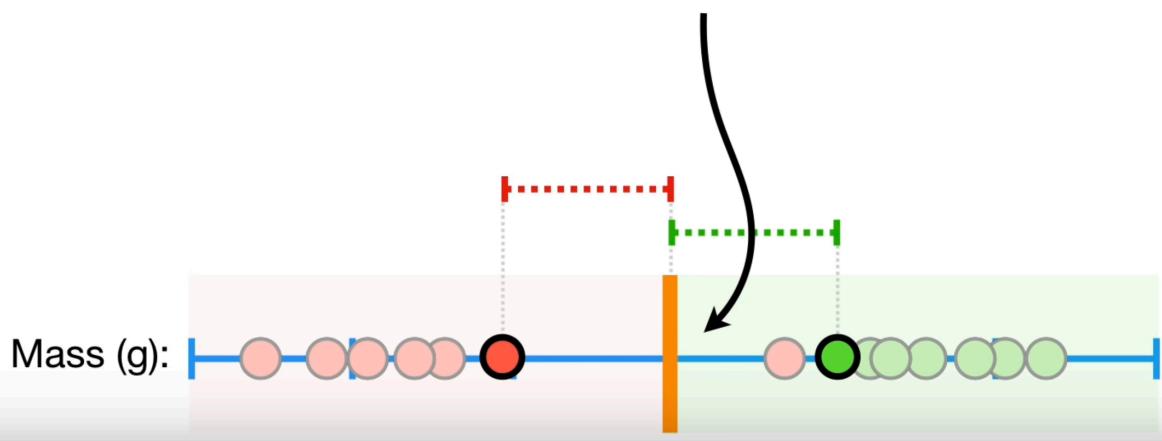
And when the data are 3-Dimensional, the **Support Vector Classifier** is a 2-Dimensional *plane* in a 3-Dimensional space.



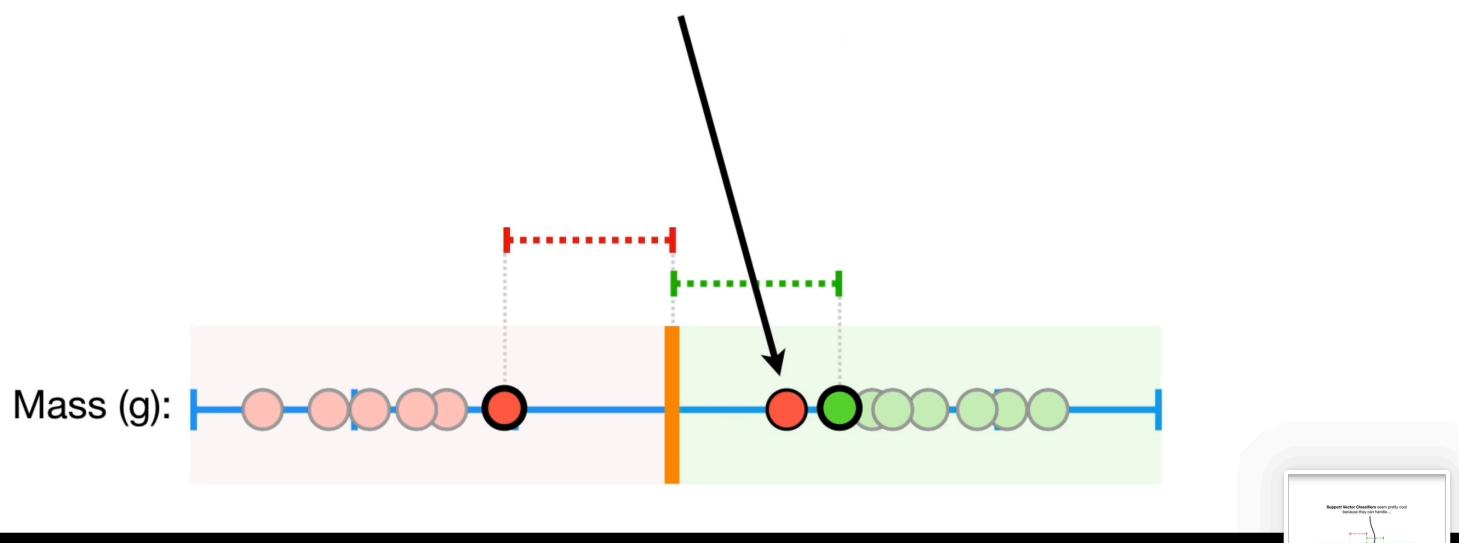
And when the data are 3-Dimensional, the Support Vector Classifier is a 2-Dimensional *plane* in a 3-Dimensional space.

psst! In mathematical jargon, a plane is a “flat affine 2-Dimensional subspace”.

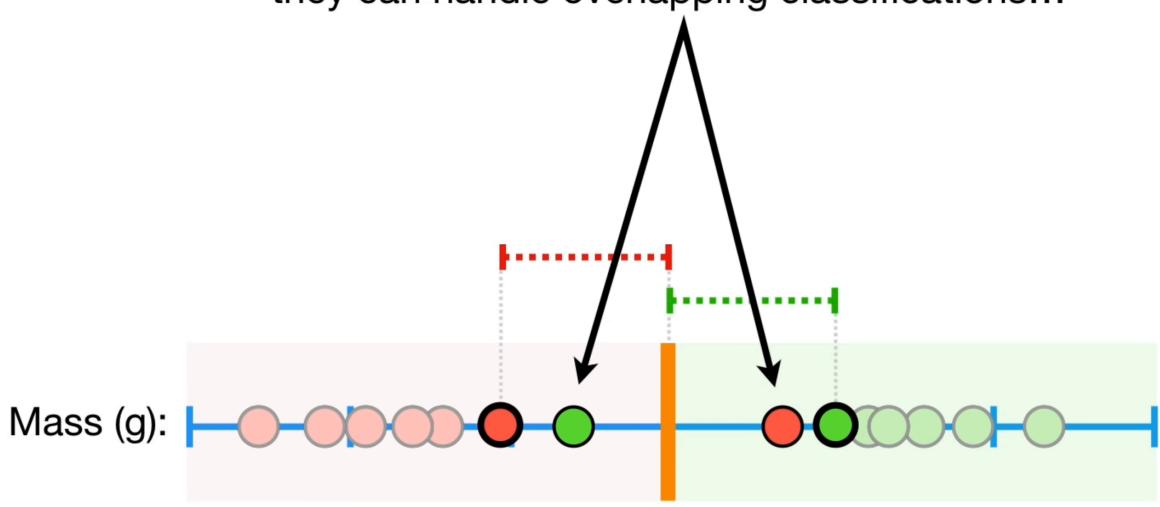
Support Vector Classifiers seem pretty cool
because they can handle...



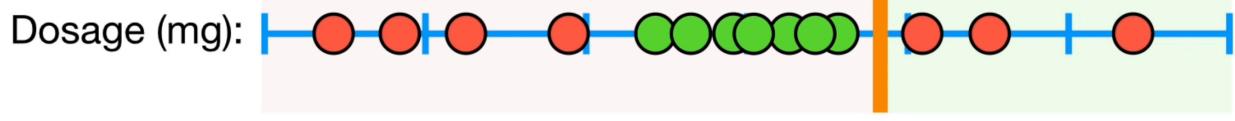
...outliers...



...and, because they allow misclassifications,
they can handle overlapping classifications...

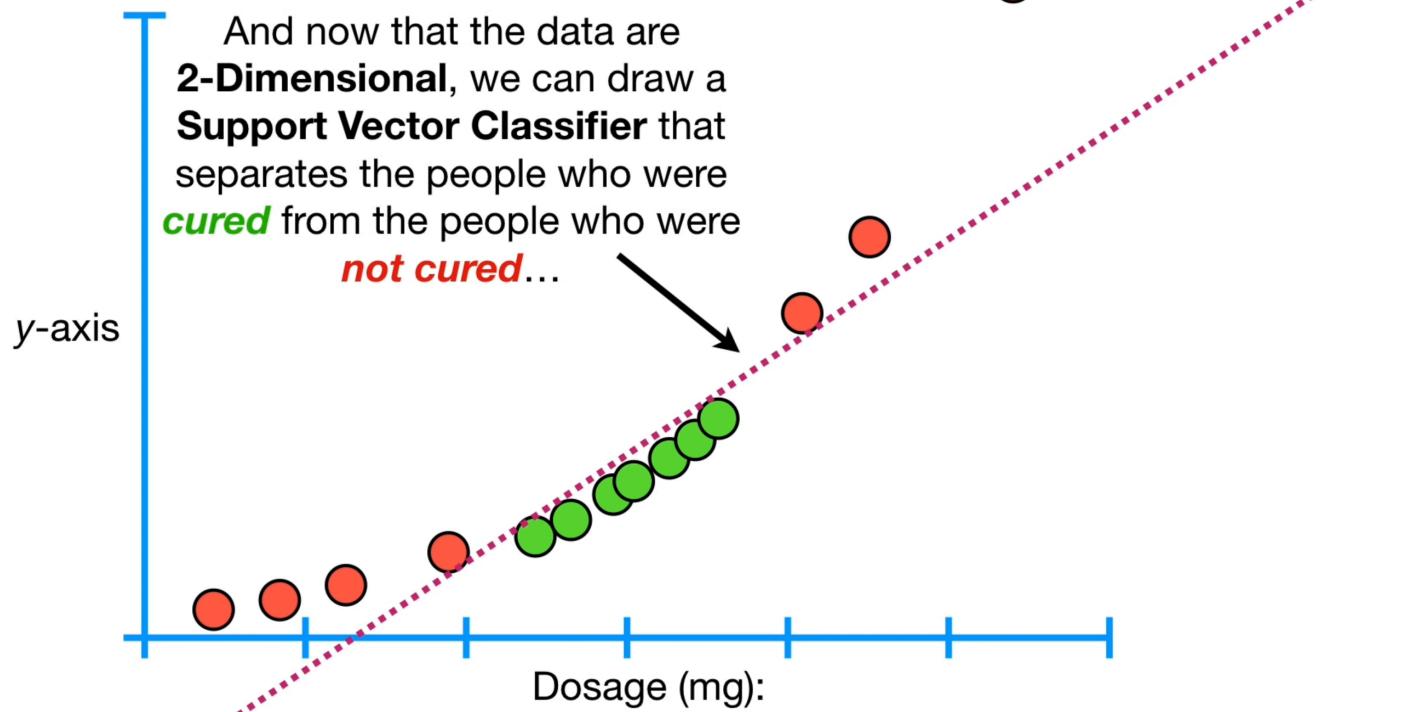


Since **Maximal Margin Classifiers** and
Support Vector Classifiers can't
handle this data, it's high time we talked
about...



Support Vector Machines!!!

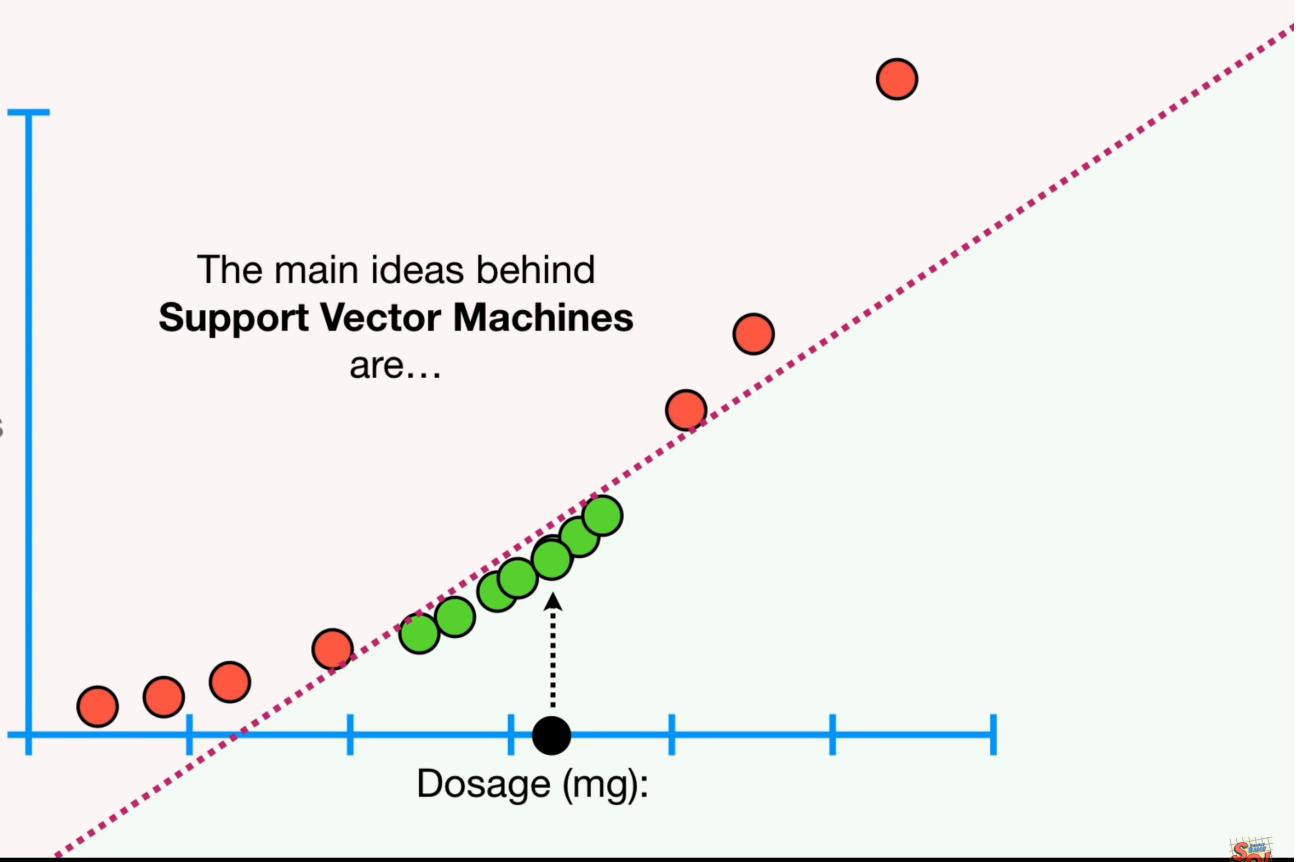
Dosage (mg): | ● ● ● | ● ● ● ● ● | ● ● | ● |

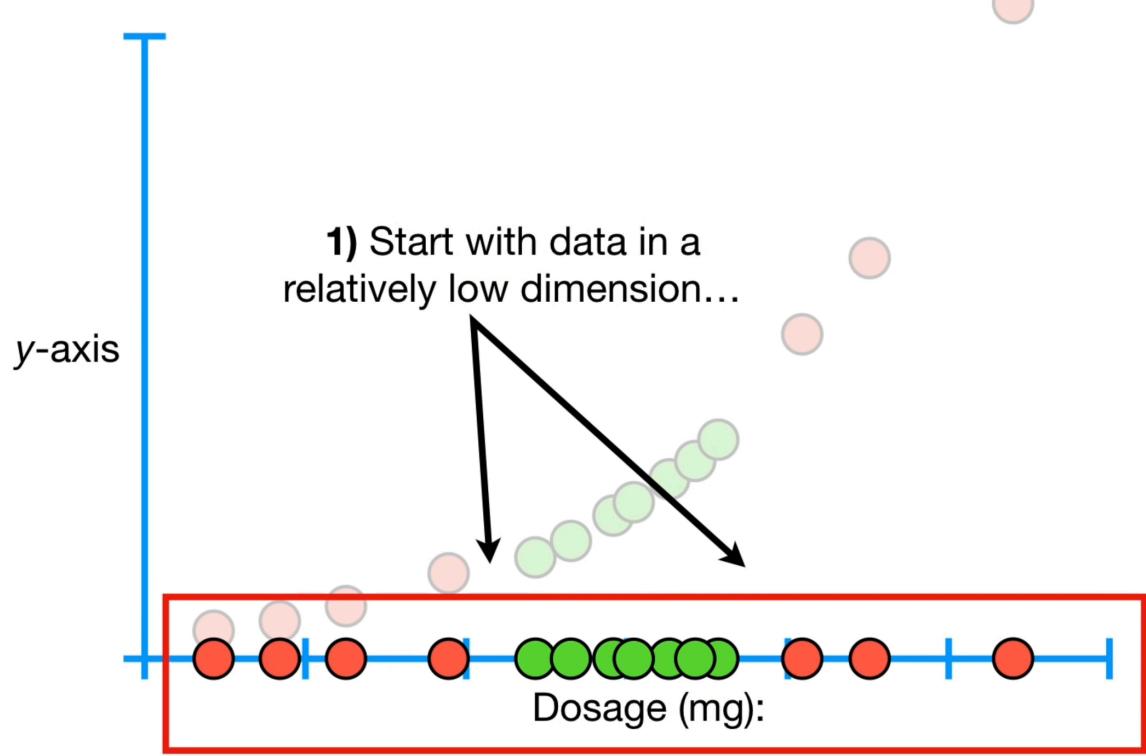


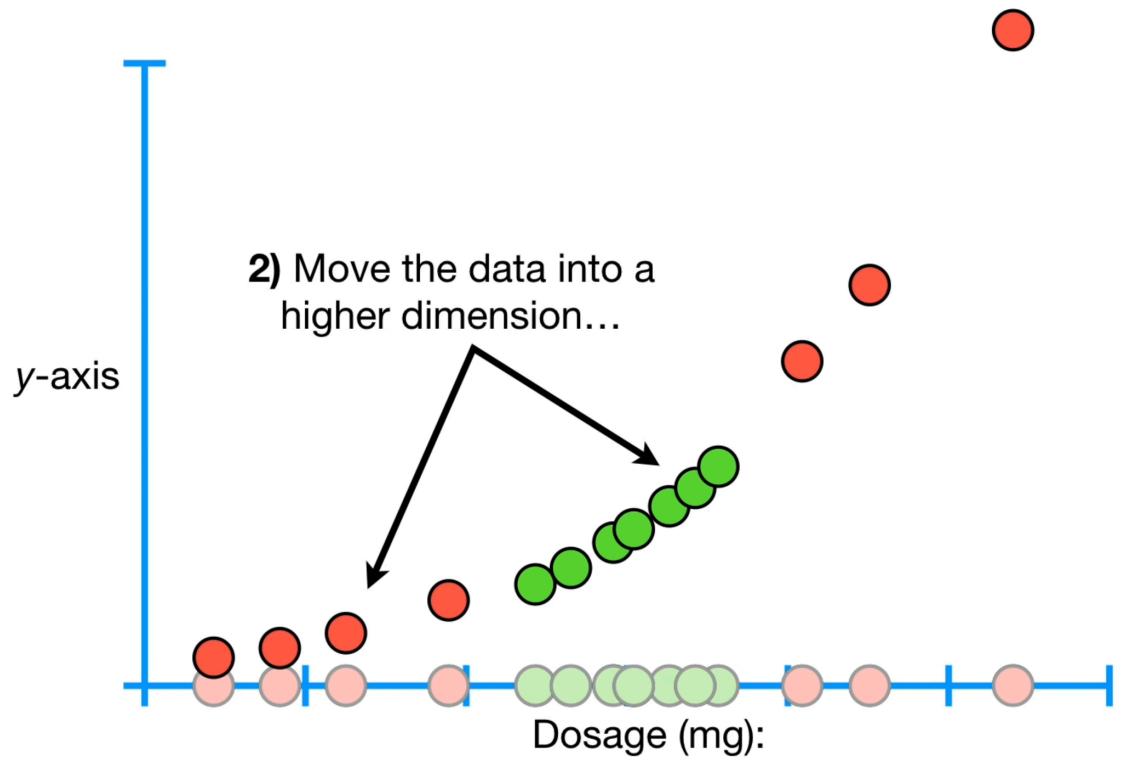
The main ideas behind
Support Vector Machines
are...

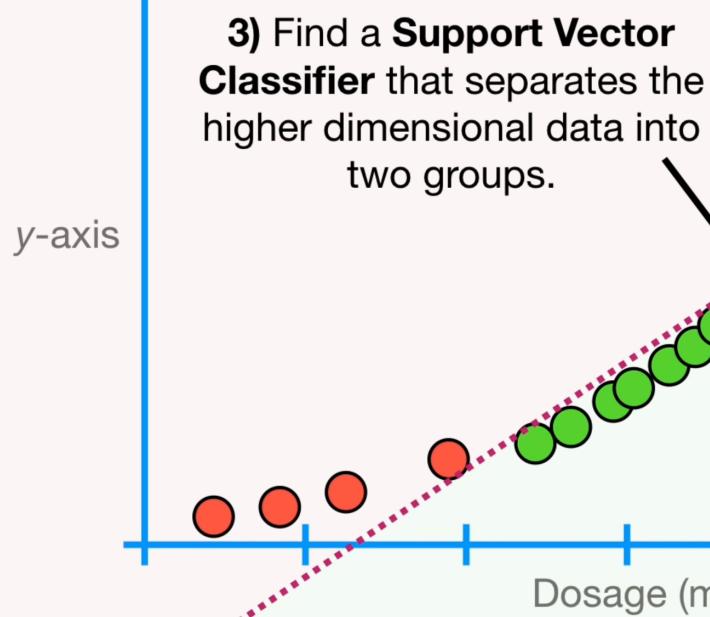
y-axis

Dosage (mg):





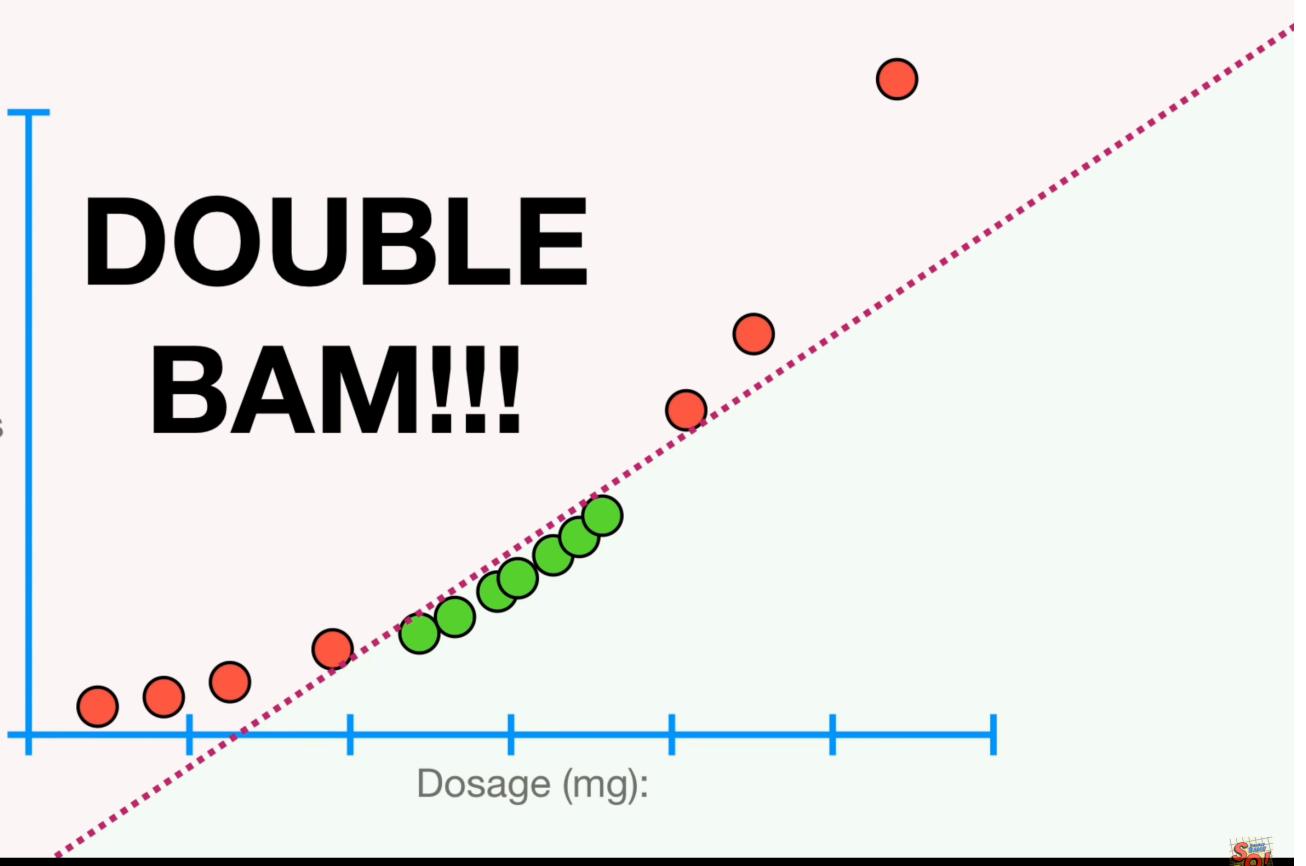


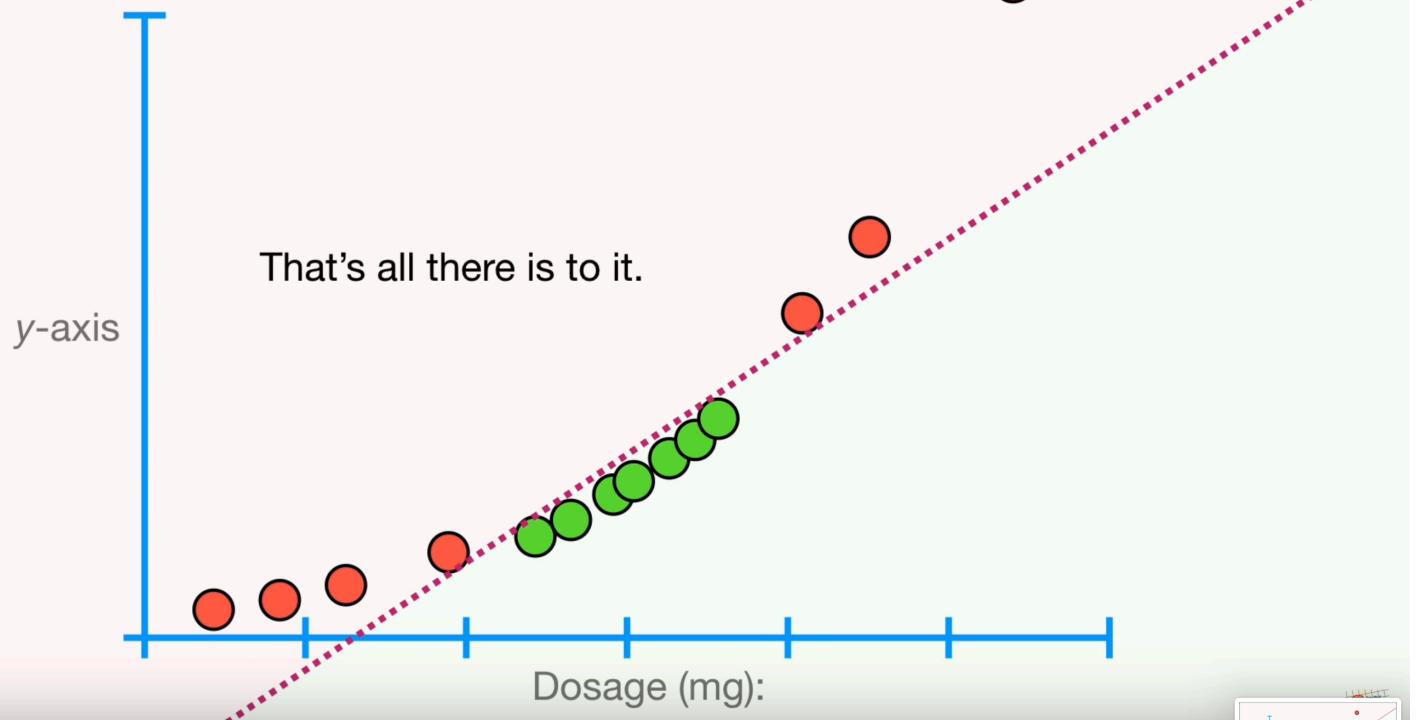


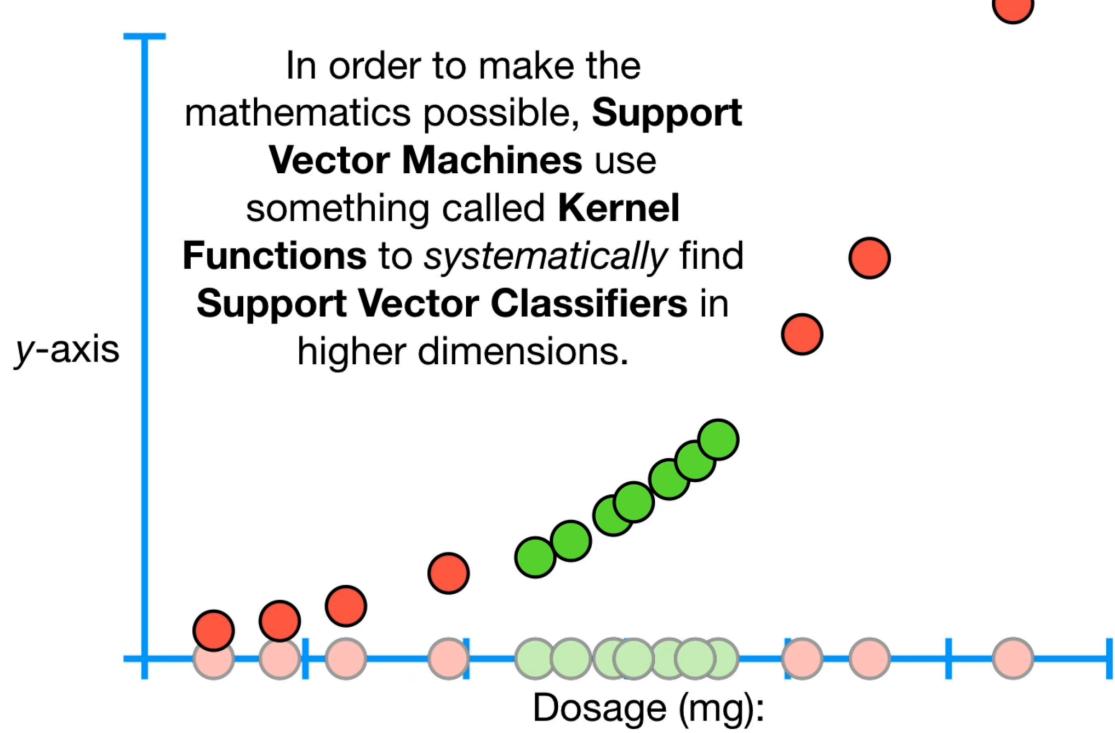
**DOUBLE
BAM!!!**

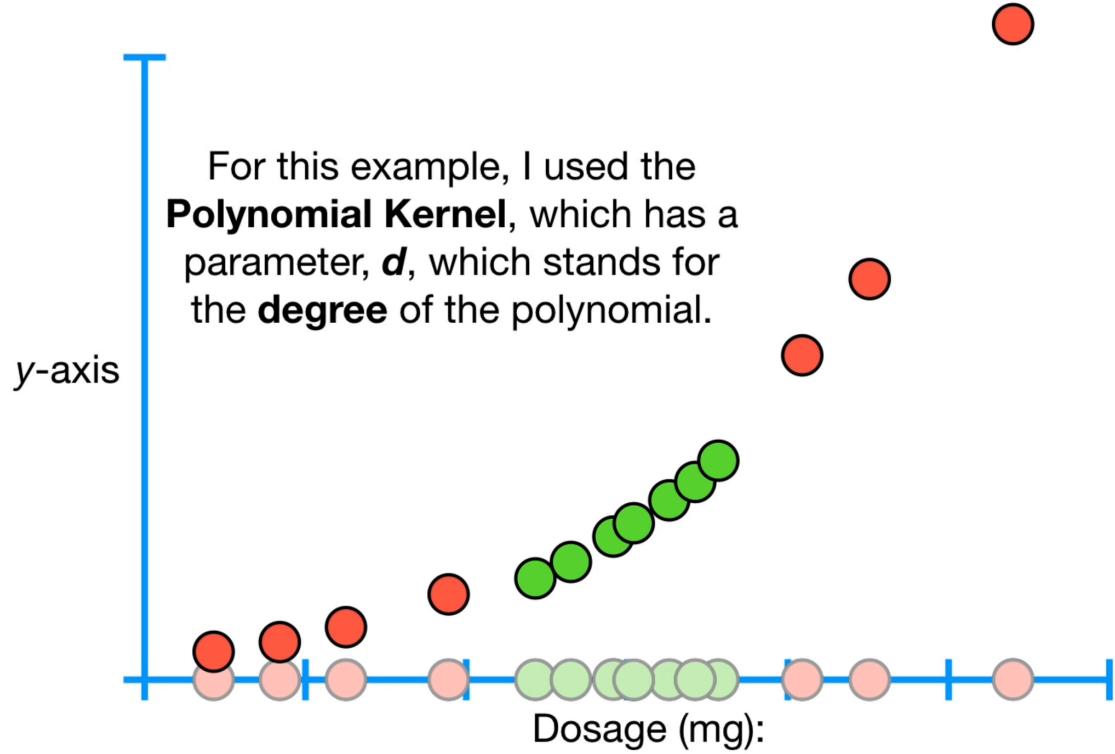
y-axis

Dosage (mg):







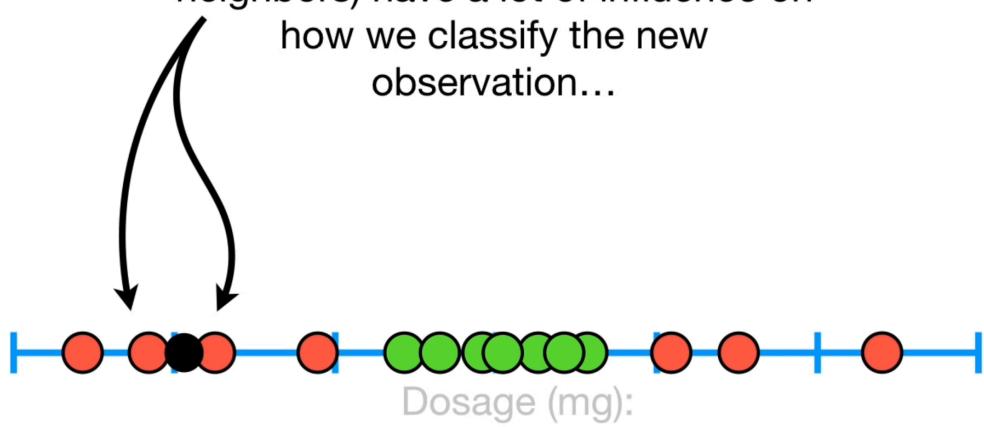


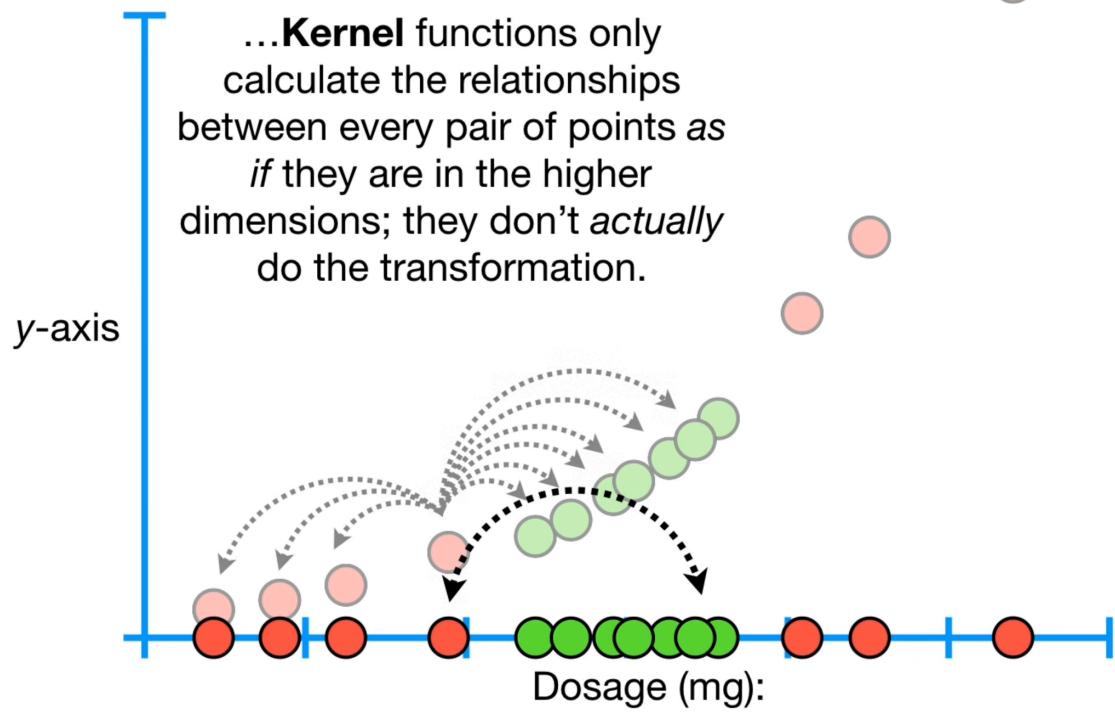
Another very commonly used **Kernel** is the **Radial Kernel**, also known as the **Radial Basis Function (RBF) Kernel**.

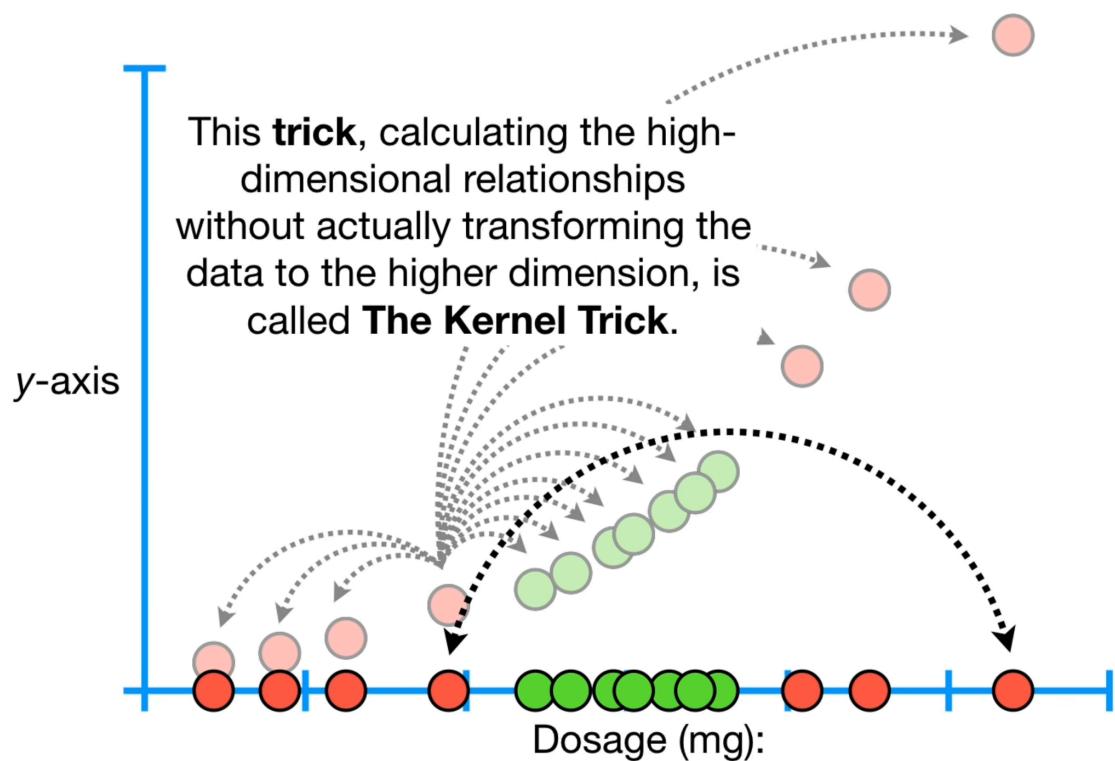
...the **Radial Kernel** behaves like a
Weighted Nearest Neighbor model.



In other words, the closest observations (aka the nearest neighbors) have a lot of influence on how we classify the new observation...







The Kernel Trick reduces the amount of computation required for **Support Vector Machines** by avoiding the math that transforms the data from low to high dimensions...

